

lixuanze_lab8

Quarto

Quarto enables you to weave together content and executable code into a finished document. To learn more about Quarto see <https://quarto.org>.

Radon

The goal of this lab is to fit this model to the radon data:

$$y_i | \alpha_{[i]} \sim N(\alpha_{[i]} + \beta x_i, \sigma_y^2), \text{ for } i=1, 2, \dots, n$$

$$\alpha_{[j]} \sim N(\gamma_0 + \gamma_1 u_j, \sigma_{\alpha}^2), \text{ for } j=1, 2, \dots, J$$

i.e. varying intercepts, fixed slope on floor. I want you to

- reproduce the graph on slide 53
- plot samples from the posterior predictive distribution for a new household in county 2 with basement level measurement, compared to samples from the posterior distribution of the mean county effect in county 2 (i.e., a graph similar to slide 45).

Here's code to get the data into a useful format:

```
library(tidyverse)
```

— Attaching core tidyverse packages — tidyverse 2.0.0 —

```
✓ dplyr      1.1.4    ✓ readr      2.1.5
✓ forcats    1.0.0    ✓ stringr    1.5.1
✓ ggplot2    3.5.0    ✓ tibble     3.2.1
✓ lubridate  1.9.3    ✓ tidyr      1.3.1
✓ purrr      1.0.2
```

— Conflicts — tidyverse_conflicts() —

```
* dplyr::filter() masks stats::filter()
```

```
* dplyr::lag()     masks stats::lag()
```

```
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
# house level data
d <- read.table(url("http://www.stat.columbia.edu/~gelman/arm/examples/radon/srrs2.dat"),

# deal with zeros, select what we want, make a fips variable to match on
d <- d |>
  mutate(activity = ifelse(activity==0, 0.1, activity)) |>
  mutate(fips = stfips * 1000 + cntyfips) |>
```

```
dplyr::select(fips, state, county, floor, activity)

# county level data
cty <- read.table(url("http://www.stat.columbia.edu/~gelman/arm/examples/radon/cty.dat"),
cty <- cty |> mutate(fips = 1000 * stfips + ctfips) |> dplyr::select(fips, Uppm)

# filter to just be minnesota, join them and then select the variables of interest.
dmn <- d |>
  filter(state=="MN") |>
  dplyr::select(fips, county, floor, activity) |>
  left_join(cty)
```

Joining with `by = join_by(fips)`

Warning in left_join(dplyr::select(filter(d, state == "MN"), fips, county, : Detected an unexpected many-to-many relationship between `x` and `y`.

```
i Row 102 of `x` matches multiple rows in `y`.
i Row 1327 of `y` matches multiple rows in `x`.
i If a many-to-many relationship is expected, set `relationship =
  "many-to-many"` to silence this warning.
```

```
head(dmn)
```

	fips	county	floor	activity	Uppm
1	27001 AITKIN		1	2.2	0.502054
2	27001 AITKIN		0	2.2	0.502054
3	27001 AITKIN		0	2.9	0.502054
4	27001 AITKIN		0	1.0	0.502054
5	27003 ANOKA		0	3.1	0.428565
6	27003 ANOKA		0	2.5	0.428565

Note, in the model:

- y_i is $\log(\text{activity})$
- x_i is floor
- u_i is $\log(\text{Uppm})$

Suggested steps

1. write Stan model (note, you will need samples from post pred distribution, either do in Stan or later in R)

Answer: Please see radon.stan file.

2. Get data in stan format

```
library(rstan)
```

Loading required package: StanHeaders

rstan version 2.32.6 (Stan version 2.32.2)

For execution on a local, multicore CPU with excess RAM we recommend calling `options(mc.cores = parallel::detectCores())`.

To avoid recompilation of unchanged Stan programs, we recommend calling `rstan_options(auto_write = TRUE)`

For within-chain threading using ``reduce_sum()`` or ``map_rect()`` Stan functions, change ``threads_per_chain`` option:

```
rstan_options(threads_per_chain = 1)
```

Attaching package: 'rstan'

The following object is masked from 'package:tidyr':

extract

```
# data transformations
N_obs <- nrow(dmn)
u <- log(dmn |> group_by(county) |> slice(1) |> select(Uppm) |> pull())
```

Adding missing grouping variables: ``county``

```
x <- dmn$floor
y <- log(dmn$activity)
county <- as.numeric(as.factor(dmn$county))
J <- length(unique(dmn$county))

stan_data <- list(y = y, x = x, u = u, N = N_obs, J = J, county = county)
```

3. Run the model

```
hirechy_model <- stan(data = stan_data, file = "radon.stan")
```

Trying to compile a simple C file

Running `/Library/Frameworks/R.framework/Resources/bin/R CMD SHLIB foo.c`

using C compiler: 'Apple clang version 15.0.0 (clang-1500.3.9.4)'

using SDK: 'MacOSX14.4.sdk'

```
clang -arch arm64 -I"/Library/Frameworks/R.framework/Resources/include" -DNDEBUG -
I"/Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/library/Rcpp/include/" -
I"/Library/Frameworks/R.framework/Versions/4.3-
arm64/Resources/library/RcppEigen/include/" -
I"/Library/Frameworks/R.framework/Versions/4.3-
arm64/Resources/library/RcppEigen/include/unsupported" -
I"/Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/library/BH/include" -
I"/Library/Frameworks/R.framework/Versions/4.3-
arm64/Resources/library/StanHeaders/include/src/" -
I"/Library/Frameworks/R.framework/Versions/4.3-
```

```
arm64/Resources/library/StanHeaders/include/" -
I"/Library/Frameworks/R.framework/Versions/4.3-
arm64/Resources/library/RcppParallel/include/" -
I"/Library/Frameworks/R.framework/Versions/4.3-arm64/Resources/library/rstan/include" -
DEIGEN_NO_DEBUG -DBOOST_DISABLE_ASSERTS -DBOOST_PENDING_INTEGER_LOG2_HPP -
DSTAN_THREADS -DUSE_STANC3 -DSTRICT_R_HEADERS -DBOOST_PHOENIX_NO_VARIADIC_EXPRESSION -
D_HAS_AUTO_PTR_ETC=0 -include '/Library/Frameworks/R.framework/Versions/4.3-
arm64/Resources/library/StanHeaders/include/stan/math/prim/fun/Eigen.hpp' -D_REENTRANT -
DRCPP_PARALLEL_USE_TBB=1 -I/opt/R/arm64/include -fPIC -falign-functions=64 -Wall -g
-O2 -c foo.c -o foo.o
```

In file included from <built-in>:1:

In file included from /Library/Frameworks/R.framework/Versions/4.3-

arm64/Resources/library/StanHeaders/include/stan/math/prim/fun/Eigen.hpp:22:

In file included from /Library/Frameworks/R.framework/Versions/4.3-

arm64/Resources/library/RcppEigen/include/Eigen/Dense:1:

In file included from /Library/Frameworks/R.framework/Versions/4.3-

arm64/Resources/library/RcppEigen/include/Eigen/Core:19:

/Library/Frameworks/R.framework/Versions/4.3-

arm64/Resources/library/RcppEigen/include/Eigen/src/Core/util/Macros.h:679:10: fatal
error: 'cmath' file not found

```
#include <cmath>
```

```
    ^~~~~~
```

1 error generated.

make: *** [foo.o] Error 1

SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 9.1e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.91 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 2000 [0%] (Warmup)

Chain 1: Iteration: 200 / 2000 [10%] (Warmup)

Chain 1: Iteration: 400 / 2000 [20%] (Warmup)

Chain 1: Iteration: 600 / 2000 [30%] (Warmup)

Chain 1: Iteration: 800 / 2000 [40%] (Warmup)

Chain 1: Iteration: 1000 / 2000 [50%] (Warmup)

Chain 1: Iteration: 1001 / 2000 [50%] (Sampling)

Chain 1: Iteration: 1200 / 2000 [60%] (Sampling)

Chain 1: Iteration: 1400 / 2000 [70%] (Sampling)

Chain 1: Iteration: 1600 / 2000 [80%] (Sampling)

Chain 1: Iteration: 1800 / 2000 [90%] (Sampling)

Chain 1: Iteration: 2000 / 2000 [100%] (Sampling)

Chain 1:

Chain 1: Elapsed Time: 0.645 seconds (Warm-up)

Chain 1: 0.603 seconds (Sampling)

Chain 1: 1.248 seconds (Total)

Chain 1:

SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 2).

Chain 2:

Chain 2: Gradient evaluation took 2.1e-05 seconds

Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.21 seconds.

Chain 2: Adjust your expectations accordingly!

Chain 2:

Chain 2:

Chain 2: Iteration: 1 / 2000 [0%] (Warmup)

Chain 2: Iteration: 200 / 2000 [10%] (Warmup)

Chain 2: Iteration: 400 / 2000 [20%] (Warmup)

Chain 2: Iteration: 600 / 2000 [30%] (Warmup)

Chain 2: Iteration: 800 / 2000 [40%] (Warmup)

Chain 2: Iteration: 1000 / 2000 [50%] (Warmup)

Chain 2: Iteration: 1001 / 2000 [50%] (Sampling)

Chain 2: Iteration: 1200 / 2000 [60%] (Sampling)

Chain 2: Iteration: 1400 / 2000 [70%] (Sampling)

Chain 2: Iteration: 1600 / 2000 [80%] (Sampling)

Chain 2: Iteration: 1800 / 2000 [90%] (Sampling)

Chain 2: Iteration: 2000 / 2000 [100%] (Sampling)

Chain 2:

Chain 2: Elapsed Time: 0.39 seconds (Warm-up)

Chain 2: 0.342 seconds (Sampling)

Chain 2: 0.732 seconds (Total)

Chain 2:

SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 3).

Chain 3:

Chain 3: Gradient evaluation took 2.8e-05 seconds

Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.28 seconds.

Chain 3: Adjust your expectations accordingly!

Chain 3:

Chain 3:

Chain 3: Iteration: 1 / 2000 [0%] (Warmup)

Chain 3: Iteration: 200 / 2000 [10%] (Warmup)

Chain 3: Iteration: 400 / 2000 [20%] (Warmup)

Chain 3: Iteration: 600 / 2000 [30%] (Warmup)

Chain 3: Iteration: 800 / 2000 [40%] (Warmup)

Chain 3: Iteration: 1000 / 2000 [50%] (Warmup)

Chain 3: Iteration: 1001 / 2000 [50%] (Sampling)

Chain 3: Iteration: 1200 / 2000 [60%] (Sampling)

Chain 3: Iteration: 1400 / 2000 [70%] (Sampling)

Chain 3: Iteration: 1600 / 2000 [80%] (Sampling)

Chain 3: Iteration: 1800 / 2000 [90%] (Sampling)

Chain 3: Iteration: 2000 / 2000 [100%] (Sampling)

Chain 3:

Chain 3: Elapsed Time: 0.393 seconds (Warm-up)

Chain 3: 0.314 seconds (Sampling)

Chain 3: 0.707 seconds (Total)

Chain 3:

SAMPLING FOR MODEL 'anon_model' NOW (CHAIN 4).

Chain 4:

Chain 4: Gradient evaluation took 2.9e-05 seconds

Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.29 seconds.

Chain 4: Adjust your expectations accordingly!

Chain 4:

Chain 4:

Chain 4: Iteration: 1 / 2000 [0%] (Warmup)

Chain 4: Iteration: 200 / 2000 [10%] (Warmup)

Chain 4: Iteration: 400 / 2000 [20%] (Warmup)

Chain 4: Iteration: 600 / 2000 [30%] (Warmup)

Chain 4: Iteration: 800 / 2000 [40%] (Warmup)

Chain 4: Iteration: 1000 / 2000 [50%] (Warmup)

Chain 4: Iteration: 1001 / 2000 [50%] (Sampling)

Chain 4: Iteration: 1200 / 2000 [60%] (Sampling)

Chain 4: Iteration: 1400 / 2000 [70%] (Sampling)

Chain 4: Iteration: 1600 / 2000 [80%] (Sampling)

Chain 4: Iteration: 1800 / 2000 [90%] (Sampling)

Chain 4: Iteration: 2000 / 2000 [100%] (Sampling)

Chain 4:

Chain 4: Elapsed Time: 0.475 seconds (Warm-up)

Chain 4: 0.514 seconds (Sampling)

Chain 4: 0.989 seconds (Total)

Chain 4:

Warning: There were 2 chains where the estimated Bayesian Fraction of Missing Information was low. See

<https://mc-stan.org/misc/warnings.html#bfmi-low>

Warning: Examine the pairs() plot to diagnose sampling problems

Warning: Bulk Effective Samples Size (ESS) is too low, indicating posterior means and medians may be unreliable.

Running the chains for more iterations may help. See

<https://mc-stan.org/misc/warnings.html#bulk-ess>

Warning: Tail Effective Samples Size (ESS) is too low, indicating posterior variances and tail quantiles may be unreliable.

Running the chains for more iterations may help. See

<https://mc-stan.org/misc/warnings.html#tail-ess>

4. For α plot, get median estimates of α 's, and the 2.5th and 97.5th percentiles. Also get the median (mean fine, easier to pull from summary) of the gamma0 and gamma1. You can then use ``geom_abline()`` to plot mean regression line.

```
extracted_samples <- rstan::extract(hirechy_model)
names(extracted_samples)
```


```
[1] "alpha"      "beta"      "gamma0"    "gamma1"    "sigma_y"
[6] "sigma_alpha" "lp__"
```

```
alpha_ <- apply(extracted_samples[["alpha"]], 2, median)
alpha_lower <- apply(extracted_samples[["alpha"]], 2, quantile, 0.025)
```

```
alpha_upper <- apply(extracted_samples[["alpha"]], 2, quantile, 0.975)

alpha_df <- tibble(county = 1:J, median = alpha_, lower = alpha_lower, upper = alpha_upper)
gamma0_ <- median(extracted_samples[["gamma0"]])
gamma1_ <- median(extracted_samples[["gamma1"]])

ggplot(alpha_df, aes(u, median)) +
  geom_point()+
  geom_errorbar(aes(ymin = lower, ymax = upper))+
  geom_abline(intercept = gamma0_, slope = gamma1_)
```



5. For the predicted y plot, you will need your posterior predictive samples for y 's and then just use `geom_density()`

```
alpha_2 <- extracted_samples[["alpha"]][,2]
sigma_y <- extracted_samples[["sigma_y"]]
y_replicated <- rnorm(alpha_2, sigma_y)

tibble(alpha = alpha_2, y = y_replicated) |>
  ggplot(aes(y)) +
  geom_density(aes(fill = "Predicted Y"), alpha = 0.6)+
  geom_density(aes(alpha, fill = "Alpha"), alpha = 0.6)
```

